

VE572 — Methods and Tools for Big Data

Midterm Rubric — Summer 2019

TA: Yihao and Yanjun (UM-JI)

This document should be kept secret.

Exercise 1 — Big data or not big data? [20 marks]

This part will count at most 20 marks in total.

1. What is the data size? [2.5+5=7.5 marks]

1. 1G-10G: Spark without Hadoop
2. 10G-100G: Mapreduce with Hadoop
3. 100G+: Spark with Hadoop

2. What is the data type? [2.5+5=7.5 marks]

1. Real time: Spark
2. Batch: Mapreduce

3. What does the company want to do with the data? [2.5+5=7.5 marks]

1. Search: Drill
2. Implement algorithms: Spark

4. Other reasonable answers [5 marks]

Exercise 2 — MapReduce [40 marks]

1. Determine all the FOFs in the following toy example. $[(2/3)*15=10]$ marks]

- 1 Ali Ben 1
- 2 Ali Col 1
- 3 Ali Gil 2
- 4 Ben Dan 1
- 5 Ben Eve 1
- 6 Ben Fin 1
- 7 Col Dan 1
- 8 Col Gil 1
- 9 Col Han 1
- 10 Dan Eve 1
- 11 Dan Han 2

```

12 Eve Fin 1
13 Eve Gil 1
14 Fin Gil 1
15 Fin Han 1

```

2. Write the Hadoop pseudocode for the first MapReduce Job. Assume a simple input text file with a list of names on each line, the user as first field followed by all his friends. For the output we expect a simple text file where each line is composed of a user and a FOF followed by the number of friends they have in common. [15 marks]

- The first MapReduce [5 marks]
- The second MapReduce [5 marks]
- The overall design [5 marks]

```

1 package com.ve572.e1;
2
3 import org.apache.commons.text.StringTokenizer;
4 import org.apache.hadoop.conf.Configuration;
5 import org.apache.hadoop.fs.FileSystem;
6 import org.apache.hadoop.fs.Path;
7 import org.apache.hadoop.io.IntWritable;
8 import org.apache.hadoop.io.Text;
9 import org.apache.hadoop.mapreduce.Mapper;
10 import org.apache.hadoop.mapreduce.Job;
11 import org.apache.hadoop.mapreduce.Reducer;
12 import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
13 import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
14
15 import java.io.IOException;
16 import java.util.ArrayList;
17
18 public class FindFOF {
19     private static String formPair(final String a, final String b) {
20         if (a.compareTo(b) < 0) return a + "," + b;
21         return b + "," + a;
22     }
23
24     public static class Map1 extends Mapper<Object, Text, Text, Text> {
25         private Text resultKey = new Text();
26         private Text resultValue = new Text();
27
28         public void map(Object key, Text value, Context context) throws IOException,
29             ↪ InterruptedException {
30             StringTokenizer tokenizer = new StringTokenizer(value.toString(), ",");

```

```

30         resultKey.set(tokenizer.nextToken());
31         while (tokenizer.hasNext()) {
32             resultValue.set(tokenizer.nextToken());
33             context.write(resultKey, resultValue);
34         }
35     }
36 }
37
38
39 public static class Reduce1 extends Reducer<Text, Text, Text, IntWritable> {
40     private Text resultKey = new Text();
41     private IntWritable resultValue = new IntWritable();
42
43     public void reduce(Text key, Iterable<Text> values, Context context) throws
44     ↪ IOException, InterruptedException {
45         ArrayList<String> arrayList = new ArrayList<>();
46         resultValue.set(0);
47         for (Text val : values) {
48             resultKey.set(formPair(key.toString(), val.toString()));
49             context.write(resultKey, resultValue);
50             arrayList.add(val.toString());
51         }
52         resultValue.set(1);
53         for (int i = 0; i < arrayList.size(); i++) {
54             for (int j = i + 1; j < arrayList.size(); j++) {
55                 resultKey.set(formPair(arrayList.get(i), arrayList.get(j)));
56                 context.write(resultKey, resultValue);
57             }
58         }
59     }
60
61     public static class Map2 extends Mapper<Object, Text, Text, IntWritable> {
62         private Text resultKey = new Text();
63         private IntWritable resultValue = new IntWritable();
64
65         public void map(Object key, Text value, Context context) throws IOException,
66         ↪ InterruptedException {
67             StringTokenizer tokenizer = new StringTokenizer(value.toString(), ",");
68             resultKey.set(tokenizer.nextToken() + " " + tokenizer.nextToken());
69             resultValue.set(Integer.parseInt(tokenizer.nextToken()));
70             context.write(resultKey, resultValue);
71         }
72     }

```

```

72
73 public static class Reduce2 extends Reducer<Text, IntWritable, Text, IntWritable> {
74     private IntWritable resultValue = new IntWritable();
75
76     public void reduce(Text key, Iterable<IntWritable> values, Context context)
77         ↪ throws IOException, InterruptedException {
78         int count = 0;
79         for (IntWritable val : values) {
80             if (val.get() == 0) return;
81             count += val.get();
82         }
83         resultValue.set(count);
84         context.write(key, resultValue);
85     }
86
87     public static void main(String[] args) throws Exception {
88         Configuration conf = new Configuration();
89         conf.set("mapred.textoutputformat.separator", ",");
90         // conf.set("mapreduce.output.textoutputformat.separator", ",");
91
92         Job job1 = Job.getInstance(conf, "ve572e1ex2.2");
93         job1.setJarByClass(FindFOF.class);
94         job1.setMapperClass(Map1.class);
95         job1.setReducerClass(Reduce1.class);
96
97         job1.setMapOutputValueClass(Text.class);
98         job1.setOutputKeyClass(Text.class);
99         job1.setOutputValueClass(IntWritable.class);
100
101         FileInputFormat.addInputPath(job1, new Path("data.txt"));
102         Path outputPath1 = new Path("output-1");
103         FileSystem fileSystem = outputPath1.getFileSystem(conf);
104         if (fileSystem.exists(outputPath1)) {
105             fileSystem.delete(outputPath1, true);
106         }
107         FileOutputFormat.setOutputPath(job1, outputPath1);
108
109         boolean exitCode = job1.waitForCompletion(true);
110         if (!exitCode) System.exit(1);
111
112         conf.set("mapred.textoutputformat.separator", " ");
113
114         Job job2 = Job.getInstance(conf, "ve572e1ex2.2");

```

```

115     job2.setJarByClass(FindFOF.class);
116     job2.setMapperClass(Map2.class);
117     job2.setReducerClass(Reduce2.class);
118
119     job2.setOutputKeyClass(Text.class);
120     job2.setOutputValueClass(IntWritable.class);
121
122     FileInputFormat.addInputPath(job2, outputPath1);
123     Path outputPath2 = new Path("output-2");
124     if (fileSystem.exists(outputPath2)) {
125         fileSystem.delete(outputPath2, true);
126     }
127     FileOutputFormat.setOutputPath(job2, outputPath2);
128     exitCode = job2.waitForCompletion(true);
129
130     System.exit(exitCode ? 0 : 1);
131 }
132
133 }

```

3. Write the Hadoop pseudocode for the second MapReduce job. Assume the previous output file as input, and as output a simple text file where each line is composed of a user and all his FOF ordered with respect to the number of common friends; for each FOF also display the number of common friends. [15 marks]

- Mapper [5 marks]
- Reducer [10 marks]

```

1  package com.ve572.e1;
2
3  import org.apache.commons.lang3.tuple.ImmutablePair;
4  import org.apache.commons.lang3.tuple.Pair;
5  import org.apache.commons.text.StringTokenizer;
6  import org.apache.hadoop.conf.Configuration;
7  import org.apache.hadoop.fs.FileSystem;
8  import org.apache.hadoop.fs.Path;
9  import org.apache.hadoop.io.Text;
10 import org.apache.hadoop.mapreduce.Job;
11 import org.apache.hadoop.mapreduce.Mapper;
12 import org.apache.hadoop.mapreduce.Reducer;
13 import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
14 import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
15
16 import java.io.IOException;
17 import java.util.ArrayList;

```

```

18 import java.util.Comparator;
19 import java.util.stream.Collectors;
20
21 public class CountFOF {
22
23     public static class Map extends Mapper<Object, Text, Text, Text> {
24         private Text resultKey = new Text();
25         private Text resultValue = new Text();
26
27         public void map(Object key, Text value, Context context) throws IOException,
28             ↳ InterruptedException {
29             StringTokenizer tokenizer = new StringTokenizer(value.toString(), " ");
30             String a = tokenizer.nextToken();
31             String b = tokenizer.nextToken();
32             String count = tokenizer.nextToken();
33             resultKey.set(a);
34             resultValue.set(b + " " + count);
35             context.write(resultKey, resultValue);
36             resultKey.set(b);
37             resultValue.set(a + " " + count);
38             context.write(resultKey, resultValue);
39         }
40     }
41
42     public static class Reduce extends Reducer<Text, Text, Text, Text> {
43         private Text resultValue = new Text();
44
45         public void reduce(Text key, Iterable<Text> values, Context context) throws
46             ↳ IOException, InterruptedException {
47             ArrayList<Pair<String, Integer>> arrayList = new ArrayList<>();
48             for (Text val : values) {
49                 String[] arr = val.toString().split(" ");
50                 arrayList.add(new ImmutablePair<>(arr[0], Integer.parseInt(arr[1])));
51             }
52             arrayList.sort(Comparator.comparing(Pair<String,
53                 ↳ Integer>::getValue).reversed());
54             resultValue.set(arrayList.stream().map(
55                 x -> x.getKey() + " " +
56                 ↳ x.getValue().toString()).collect(Collectors.joining(", ")));
57             context.write(key, resultValue);
58         }
59     }
60 }

```

```

58     public static void main(String[] args) throws Exception {
59         Configuration conf = new Configuration();
60         conf.set("mapred.textoutputformat.separator", " ");
61         //     conf.set("mapreduce.output.textoutputformat.separator", ",");
62
63         Job job = Job.getInstance(conf, "ve572e1ex2.3");
64         job.setJarByClass(CountF0F.class);
65         job.setMapperClass(CountF0F.Map.class);
66         job.setReducerClass(CountF0F.Reduce.class);
67
68         job.setOutputKeyClass(Text.class);
69         job.setOutputValueClass(Text.class);
70
71         FileInputFormat.addInputPath(job, new Path("output-2"));
72         Path outputPath = new Path("output-3");
73         FileSystem fileSystem = outputPath.getFileSystem(conf);
74         if (fileSystem.exists(outputPath)) {
75             fileSystem.delete(outputPath, true);
76         }
77         FileOutputFormat.setOutputPath(job, outputPath);
78
79         boolean exitCode = job.waitForCompletion(true);
80         System.exit(exitCode ? 0 : 1);
81     }
82 }

```

Exercise 3 — Course questions [35 marks]

1. HDFS [15 marks]

a) What is the default replication level in HDFS? [1 mark]

The default replication factor is 3.

b) Parallel is often seen as more efficient than serial. When writing a file in HDFS, blocks are first sent to a DataNode which forwards them to another, which sends them to another, and so on...Why is this process not parallelised, i.e. send blocks to all the DataNodes at the same time? [4 marks]

- If the blocks are sent to all the DataNodes at the same time, it will be slow due to the limited throughput. [2 marks]
- Thus when writing a file in HDFS, the blocks are sent to a DataNode at a time. Once the blocks are written in a DataNode, it is no need to forward them to other DataNodes. [2 marks]

c) Explain how to find a file in HDFS. [3 marks]

1. Each datanode announces the blocks it has. [1 mark]
2. The namenode keeps all the information in its memory. [1 mark]
3. When a write occurs an entry is added to the edit log. [1 mark]

d) What are the namespace image and edit log? [3 marks]

- Namespace image: Storing the entire file system namespace, including the mapping of blocks to files and file system properties. [1.5 marks]
- Edit log: Recording every change that occurs to file system metadata, such as creating a new file in HDFS and changing the replication factor of a file. [1.5 marks]

e) Is it possible to have several NameNodes in a cluster? If so explain how it works. [4 marks]

- It is possible. [1 mark]
- We can use HDFS federation. [1 mark]
- Split the filesystem over several independent namenodes. Each namenode has a namespace and its own pool of blocks [1 mark]
- A namespace with a block pool is called namespace volume and a datanode is not attached to a specific namespace volume. [1 mark]

2. YARN [8 marks]

a) Explain how an application is launched and run using YARN. [3 marks]

1. Application client submit a YARN application to Resource Manager.
2. Resource Manager contacts Node Manager to launch a new container and run Application Master in it. [1 mark]
3. Application Master asks Resource Manager for allocating the resources. [1 mark]
4. Application manager gets the resources information from Resource Manager and it launches the container through other Node Manager. [1 mark]

b) Would you recommend the fair or capacity scheduler? Explain the when and why. [5 marks]

- – Fair scheduling is a method of assigning resources to applications such that all apps get, on average, an equal share of resources over time. [1 mark]
 - It is a good default for small to medium sized clusters [0.5 mark]
 - since it is more flexible and allows for jobs to consume unused resources in the cluster. [1 mark]
- – Capacity scheduler is designed to run Hadoop applications as a shared, multi-tenant cluster in an operator-friendly manner while maximizing the throughput and the utilization of the cluster. [1 mark]
 - It's generally used on large clusters with lots of different workloads with different needs [0.5 mark]

- since it can give each organization capacity guarantees.[1 mark]

3. Briefly explain the similarities and differences between MapReduce and Spark. [7 marks]

- Similarity: Both are highly scalable and can be used in a cluster. [1 mark]
- MapReduce:
 - MapReduce takes two stages to process data, Map and Reduce. It reads and writes from disk and thus slower and has high latency. [1 mark]
 - It uses replication for fault tolerance, which significantly increase the completion times for operations with a single failure. [2 marks]
- Spark:
 - Spark is lightning fast cluster computing tool. It is much faster than MapReduce and has low-latency computing. [1 mark]
 - It uses RDDs and DAG for fault tolerance. If an RDD is lost, it is easy to recompute a new one by using the original transformations. [2 marks]

4. Drill [5 marks]

a) What is Zookeeper, and why is it a core component of Drill's strategy? [2 marks]

- ZooKeeper is a centralized service for maintaining configuration information, naming, providing distributed synchronization, and providing group services. [1 mark]
- Drill uses ZooKeeper to maintain cluster membership and health-check information. [1 mark]

b) Explain what is a Drillbit and how it functions. [3 marks]

- A Drillbit is the process running on each active Drill node that coordinates, plans, and executes queries, as well as distributes query work across the cluster to maximize data locality. [1 mark]
- The Drillbit receives the query from a client. A SQL parser in the DrillBit parses the SQL and form a logical plan. [0.5 mark]
- The Drillbit sends the logical plan into a optimizer to optimize and convert the logical plan into a physical plan that describes how to execute the query. [0.5 mark]
- A parallelizer in the Drillbit transforms the physical plan into multiple phases, called major and minor fragments. [0.5 mark]
- These fragments create a multi-level execution tree that rewrites the query and executes it in parallel against the configured data sources, sending the results back to the client or application. [0.5 mark]

Exercise 4 — Simple Hadoop questions [5 marks]

1. Why is ssh needed on the master and workers? How to configure it? [1 mark]

- The master uses ssh protocol to send commands to the workers. [0.5 mark]

- The worker should add the master's public key in ssh configurations. [0.5 mark]

Which Java version is needed by Hadoop, why? [1 mark]

- Java version 8 is needed. [0.5 mark]
- Some APIs are deprecated in the new versions of Java and Hadoop hasn't altered them yet. [0.5 mark]

Why should Hadoop's home be the same across the whole cluster? [1 mark]

Because the master use the same configuration file to find the Hadoop's home on every worker.

How to use hdfs dfs command line interface to [2 marks]

(i) list a directory [1 mark]

```
hdfs dfs -ls <hdfs path>
```

(ii) upload or download a file [1 mark]

- `hdfs dfs -put <local path> <hdfs path>` [0.5 mark]
- `hdfs dfs -get <hdfs path> <local path>` [0.5 mark]